

Hotline Networking

Eric Grosse grosse@gmail.com

version 2020-02-07 latest version will be at github.com/n2vi/hotline

This document makes a concrete proposal for the network link between the Puck and ROCCS hotline devices owned by one country. We seek a small, easily reviewed network stack with minimal OS support. It should depend on little in the surrounding office environment.

In the CATALINK architecture, there is a small device called the Puck owned by a hotline Principal that holds all secrets for the end-to-end encrypted hotline. It may only be intermittently powered on, so relies for connectivity on the ROCCS, a well-connected and always-on relay server, also owned by the Principal but with somewhat less security sensitivity. The goal of this document is to describe how Puck talks to ROCCS.

Puck is expected to run a microkernel such as seL4 plus a small amount of userland software that should be reviewable by people from many backgrounds. A typical TCP/IP network stack would be larger than the entire rest of the microkernel so the CATALINK team identified early on that a simpler networking alternative is a crucial element of technical credibility for our project.

ROCCS is expected to run a larger operating system such as Linux or OpenBSD and, besides talking to Puck, be capable of handling many other communication channels such as Internet, private fiber, satellite links, and HF radio. We will endeavor to make it secure but recognize that with the necessary added complexity comes added risk. So ROCCS holds mild secrets, enough to deter most wiretapping or impersonation, but nothing crucial to hotline message content.

In CATALINK we believe we can achieve radically better simplicity because we are not trying for ultimate performance. Hardware has gotten orders of magnitude faster since the '60s era Moscow-Washington hotline that we're reconstructing and we're willing to spend some of that on simplicity. Not that performance doesn't matter; our final system needs to feel snappy. Perhaps even more important than hardware speeds, we believe an enabling factor for better security is that we have few legacy constraints.

Puck and ROCCS are imagined to live in the same office building and connected by CAT5 cable, possibly through a Gigabit Ethernet switch. We assume a low bandwidth-delay product and a reasonably uncongested network. We do not entirely trust the network, though we are willing to assume that adversaries would find it challenging to disrupt. More specifically, from a security threat perspective, we feel the need for encryption to protect against wiretap or forgery but do not feel the need for extraordinary defense against Denial-of-Service packet flooding or against physical disconnection. To the extent feasible, we would like any network outage to be diagnosable by relatively unskilled staff, who can spot unplugged or frayed cables or power failures, but not DNS misconfiguration. Initial setup may require an expert.

The [Precursor](#) device uses WiFi rather than wired ethernet but the security relevant network stack is a COM bus that is also designed for simplicity, so we're willing to switch to that if project hardware choices dictate.

It is possible that a future design will allow for multiple ROCCS (think one in the White House and one on Air Force One) exchanging state when able. For messages to get through it would be enough that the Puck be able to talk to any one of these. For the rest of this document, assume a single ROCCS.

The current state of the proposal is aimed at prototypes on common platforms but prepares for a tiny OS on future devices, and therefore we imagine a primitive UDP interface provided by the operating system and limit what we ask of it to handling the network hardware and multiplexing packets (based on UDP port number) to and from independently scheduled user-mode applications which are responsible for their own retransmission and encryption. We do not ask for DNS or routing configuration or the rest of a normal IP infrastructure.

puckfs (For historical origins, see [Plan9 rudp](#).)

Instead of general message passing, for more understandable control flow we conjecture all we need is remote procedure call. Furthermore, rather than general RPC, we think we can get by with a primitive remote filesystem, reading and writing full files. For simplicity, we decided:

- no walk, always use canonicalized pathname
- no stat, that info comes when reading directory
- no mkdir; directories are created as needed when file is written
- no symlinks, no hard links, no locking, no ACLs

The **client** calls

```
p, err = puckfs.Dial(secretFile)
data, err = p.ReadFile(path)
err = p.WriteFile(path, data)
err = p.Remove(path)
fileinfo, err = p.ReadDir(path)
err = p.Chtime(path, mtime)
err = p.Close()
```

Here `secretFile` is a local file containing JSON fields: `RemoteAddr`, `MTU`, `KeyID`, `Secret`. `KeyID` is unique to a (client,server) pair and is used in both directions. Calls are synchronous. The filesystem is single-user, so has no permissions.

The **server** process is deliberately primitive, supporting a single remote user and running as a custom local userid with access to just the one dedicated directory tree.

If we need more functionality, try files first. For printing, as an example, don't implement CUPS but `WriteFile` a PDF from the puck to a spool directory on ROCCS.

Q: Why the belt-and-suspenders construction of nonce in the implementation?

A: Crypto-random bytes are a well-accepted way to implement a nonce, so we wanted to use that for sure. To ease concerns about subverted randomness generators, it was mildly attractive to also have an independent means of assuring the nonce is unique. There are presumably other bad things cryptographically that happen if your random source isn't. So take pains to check that.

The main reason to have monotonic counters with fatal errors is as a tripwire for otherwise undetected theft of the secret. I acknowledge there may be operational pain in keeping global counters in sync, even just between puck and ROCCS, so might have to revisit whether this is treated as a fatal error but will try for now.

There is intended to be no downside to using both. As long as either mechanism is working, even if the other is controlled by an adversary, the nonce should be safe. If not, that's a critical bug.

Q: What is the user interface?

A: Since the puck hardware is still to be settled, discussion of UI is somewhat premature. For this proof-of-concept implementation, we're building on top of OpenBSD and running commands within the acme text editor. Error messages are seen directly on that screen.

TODO

We must provide a safe way to set (and reset) the shared secret between client and server. PAKE seems a possible choice, but let's get the rest settled first.

What is the best way to rigorously define all this, once we have the right design? Our ultimate goal is to apply formal methods to assess the system, so whatever can be done to make that more feasible is a good thing. Is SPIN still the right tool for protocols?

We want some **notification** channel, for example to be alerted when a directory changes. We also want something similar to **mossh**.

I thank Dave Presotto for comments on this proposal and even more for teaching me all this in the Plan 9 days.